

IMD0905 - Data Science I

Lesson #9 - Introduction to Pandas

Ivanovitch Silva
September, 2018



Agenda

- About the two core pandas types: dataframes and series
- How to select data using row and column labels
- A variety of methods for exploring data with pandas
- How to assign data using various techniques in pandas
- How to use boolean indexing with pandas for selection and assignment

Update the repository

```
git clone https://github.com/ivanovitchm/IMD0905_datascience_one.git
```

Or

```
git pull
```

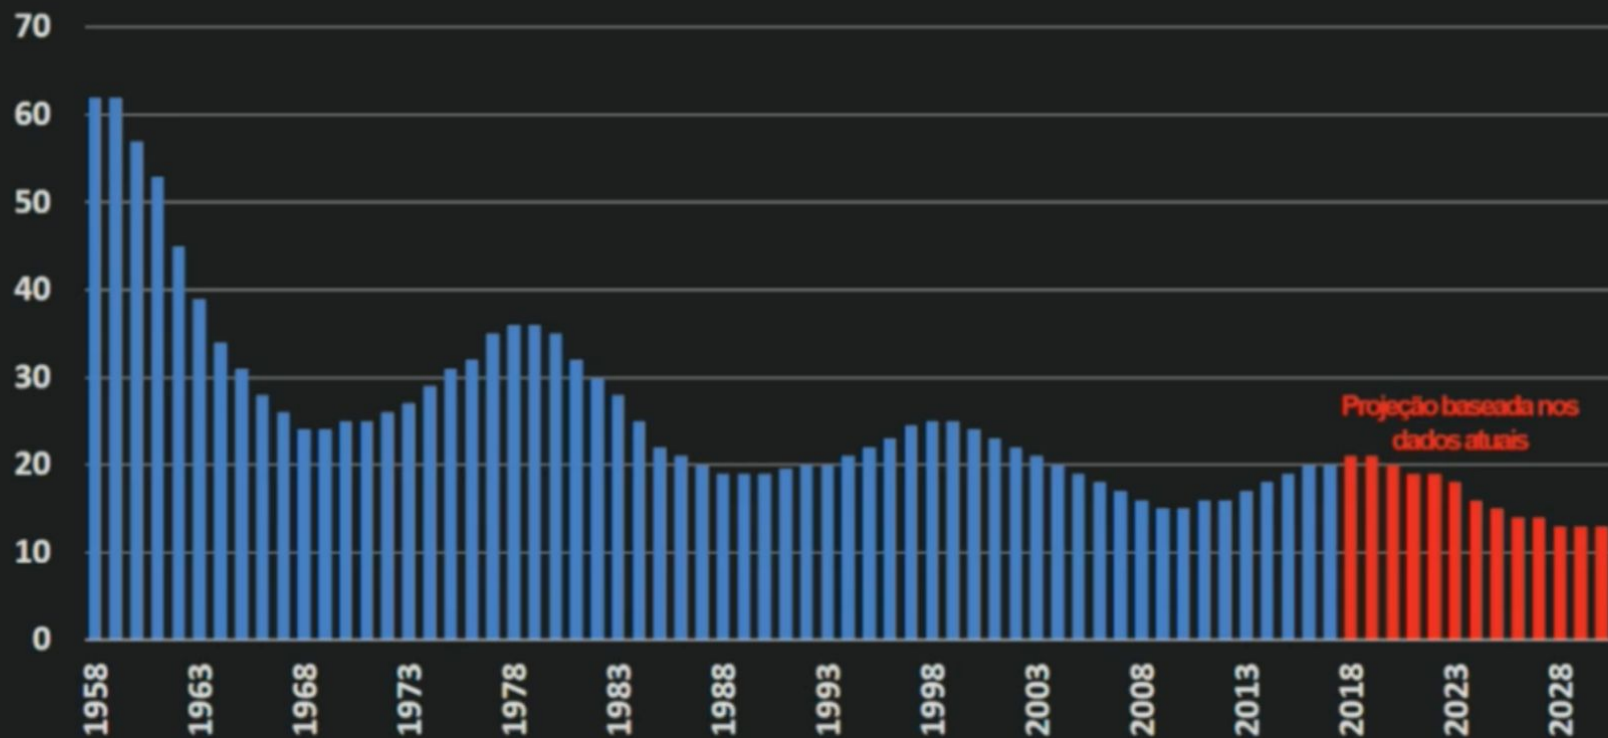
Understanding Pandas & Numpy

- Numpy
 - a. Lack support for column names
 - b. Support for only one data type per ndarray
 - c. There are lots of low level methods, however there are many common analysis patterns that don't have pre-built methods.

The pandas library provides solutions to all of these pain points and more. Pandas is not so much a replacement for NumPy as an **extension of NumPy**.



Tempo médio de permanência de uma empresa no S&P 500 (em anos)



Fonte: INNOSIGHT, Richard N. Foster, Standard & Poor's

UMA EMPRESA DO S&P 500 ESTÁ SENDO SUBSTITUÍDA A CADA DUAS SEMANAS


Richard Foster

The dataset

	rank	revenues	revenue_change	profits	assets	profit_change	ceo	industry	sector	previous_rank
Walmart	1	485873	0.8	13643.0	198825	-7.2	C. Douglas McMillon	General Merchandisers	Retailing	1
State Grid	2	315199	-4.4	9571.3	489838	-6.2	Kou Wei	Utilities	Energy	2
Sinopec Group	3	267518	-9.1	1257.9	310726	-65.0	Wang Yupu	Petroleum Refining	Energy	4
China National Petroleum	4	262573	-12.3	1867.5	585619	-73.7	Zhang Jianhua	Petroleum Refining	Energy	3
Toyota Motor	5	254694	7.7	16899.3	437575	-12.3	Akio Toyoda	Motor Vehicles and Parts	Motor Vehicles & Parts	8

```
import pandas as pd
f500 = pd.read_csv("f500.csv", index_col=0)
f500.index.name = None
```


Introducing Dataframes



Introducing Dataframes

```
# put your code here  
f500.info(,)
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 500 entries, Walmart to AutoNation  
Data columns (total 16 columns):  
rank                500 non-null int64  
revenues            500 non-null int64  
revenue_change      498 non-null float64  
profits             499 non-null float64  
assets              500 non-null int64  
profit_change       436 non-null float64  
ceo                  500 non-null object  
industry             500 non-null object  
sector              500 non-null object  
previous_rank        500 non-null int64  
country              500 non-null object  
hq_location          500 non-null object  
website              500 non-null object  
years_on_global_500_list 500 non-null int64  
employees            500 non-null int64  
total_stockholder_equity 500 non-null int64  
dtypes: float64(3), int64(7), object(6)  
memory usage: 66.4+ KB
```

f500.head()
f500.tail()

Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
f500_selection	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

```
f500_selection.loc[:, "rank"]
```

Walmart	1
State Grid	2
Sinopec Group	3
China Natural Petroleum	4
Toyota Motor	5

Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
f500_selection	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

		country	rank
f500_selection.loc[:, ["country", "rank"]]	Walmart	USA	1
	State Grid	China	2
	Sinopec Group	China	3
	China Natural Petroleum	China	4
	Toyota Motor	Japan	5

Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
f500_selection	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

		rank	revenues	profits
f500_selection.loc[:, "rank": "profits"]	Walmart	1	485873	13643.0
	State Grid	2	315199	9571.3
	Sinopec Group	3	267518	1257.9
	China Natural Petroleum	4	262573	1867.5
	Toyota Motor	5	254694	16899.3

Column selection shortcuts


Select by Label	Explicit Syntax	Common Shorthand	Other Shorthand
Single column	<code>df.loc[:, "col1"]</code>	<code>df["col1"]</code>	<code>df.col1</code>
List of columns	<code>df.loc[:, ["col1", "col7"]]</code>	<code>df[["col1", "col7"]]</code>	
Slice of columns	<code>df.loc[:, "col1":"col4"]</code>		

Lesson_09_Introduction_to_pandas.ipynb

Up to section 1.4



Selecting Items from a Series by Label



Original
Dataframe

	A	B	C	D	E
V					
W					
X					
Y					
Z					

Code

```
single_col = df["D"]
```

Result

V	
W	
X	
Y	
Z	

single_col is a
series object

	A	B	C	D	E
V					
W					
X					
Y					
Z					

```
single_row = df.head(1)
```

A	
B	
C	
D	
E	

single_row is a
series object

Original
Dataframe

	A	B	C	D	E
V					
W					
X					
Y					
Z					

Code

```
multi_cols = df[["A", "C", "D"]]
```

Result

	A	C	D
V			
W			
X			
Y			
Z			

multi_cols is a
dataframe object

	A	B	C	D	E
V					
W					
X					
Y					
Z					

```
multi_rows = df.head(3)
```

	A	B	C	D	E
V					
W					
X					

multi_rows is a
dataframe object

Dataframe vs Series

	Series	DataFrame
Dimensions	One	Two
Has 'index' axis	Yes	Yes
Has 'columns' axis	No	Yes
Number of dtypes	One	Many (one per column)

Series and Dataframe Describe Methods

```
revs = f500["revenues"]  
print(revs.describe())
```

```
count      500.000000  
mean      55416.358000  
std       45725.478963  
min       21609.000000  
25%       29003.000000  
50%       40236.000000  
75%       63926.750000  
max       485873.000000  
Name: revenues, dtype: float64
```

```
print(f500["assets"].describe())
```

```
count      5.000000e+02  
mean      2.436323e+05  
std       4.851937e+05  
min       3.717000e+03  
25%      3.658850e+04  
50%      7.326150e+04  
75%      1.805640e+05  
max      3.473238e+06  
Name: assets, dtype: float64
```


Lesson_09_Introduction_to_pandas.ipynb

Sections 1.5, 1.6, 1.7



`DataFrame.method(axis=0)`
or
`DataFrame.method(axis="index")`


Calculates along the **row** axis



Calculates result for
each **column**.

`DataFrame.method(axis=1)`
or
`DataFrame.method(axis="column")`

Calculates along the **column** axis



Calculates result for
each **row**.

More data exploration methods

```
medians = f500[["revenues", "profits"]].median(axis=0)
# we could also use .median(axis="index")
print(medians)
```

```
revenues    40236.0
profits      1761.6
dtype: float64
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	485873
State Grid	2	315199
Sinopec Group	3	267518
China National Petroleum	4	262573
Toyota Motor	5	254694

```
>>> top5_rank_revenue["revenues"] = 0
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	0
State Grid	2	0
Sinopec Group	3	0
China National Petroleum	4	0
Toyota Motor	5	0

Assignment with Pandas

Assignment with Pandas

```
>>> top5_rank_revenue.loc["Sinopec Group", "revenues"] = 999
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	0
State Grid	2	0
Sinopec Group	3	999
China National Petroleum	4	0
Toyota Motor	5	0

Add a new column

```
>>> top5_rank_revenue["year_founded"] = 0
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues	year_founded
Walmart	1	0	0
State Grid	2	0	0
Sinopec Group	3	999	0
China National Petroleum	4	0	0
Toyota Motor	5	0	0

Add a new row

```
>>> top5_rank_revenue.loc["My New Company"] = 555
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues	year_founded
Walmart	1	0	0
State Grid	2	0	0
Sinopec Group	3	999	0
China National Petroleum	4	0	0
Toyota Motor	5	0	0
My New Company	555	555	555

Using boolean indexing with pandas objects

w	2
x	4
y	6
z	8

pandas series



w	2	< 5
x	4	< 5
y	6	< 5
z	8	< 5

Vectorized
boolean operation



w	True
x	True
y	False
z	False

Boolean
pandas series

	A	B
w	2	3
x	4	6
y	6	9
z	8	12

pandas DataFrame



	A	B	
w	2	3	< 5
x	4	6	< 5
y	6	9	< 5
z	8	12	< 5

Vectorized
boolean operation



	A	B
w	True	True
x	True	False
y	False	False
z	False	False

Boolean
pandas DataFrame

Using boolean indexing with pandas objects

```
result = df.loc[num_bool, "name"]
```

		name		
False	w	Kylie	12	
True	→ x	Rahul	8	→ x
False	y	Michael	5	z
True	→ z	Sarah	8	↗
				result
				Rahul
				Sarah

```
result = df[num_bool]
```

		name	num	
False	w	Kylie	12	
True	→ x	Rahul	8	→ x
False	y	Michael	5	z
True	→ z	Sarah	8	↗
				result
				Rahul
				Sarah

Using boolean arrays to assign values

```
f500.loc[f500["sector"] == "Motor Vehicles & Parts", "sector"] = "Motor Vehicles and Parts"
```

Challenge

Finding top performers by country

```
>>> top_3_countries = f500["country"].value_counts().head(3)
```

```
>>> print(top_3_countries)
```

```
USA      132
```

```
China    109
```

```
Japan     51
```

```
Name: country, dtype: int64
```

Lesson_09_Introduction_to_pandas.ipynb

